



Embedded Linux  
Conference  
North America

- Tracing resource-constrained embedded systems using eBPF

**Adrian Ratiu, Collabora Ltd**



# Agenda

- About me
- Embedded / IoT woes
- How does eBPF fit in?
- Quick eBPF / BCC introduction, benefits
- Approaches to eBPF on embedded devices
- Trade-offs, specific projects pros/cons
- Ways forwards



## About me.

I enjoy



working in a company of awesome FOSS-oriented  
people at Collabora

work with companies who “get it” when using FOSS

work to help companies “get it” and be successful



## I also really enjoy

Taking systems apart and modifying them

Projects like OpenEmbedded/Yocto, Buildroot/OpenWRT

Always looking for new tech to improve development  
and debugging of embedded devices

Learning about eBPF (just a user, not an expert)

A strong dislike of locked-down devices /  
that lock owner usage without very good reasons

# Embedded and the IoT

- “Smart” devices everywhere
- Increasingly powerful, complex, connected hardware
- Much more capable than default software installations allow
- Software complexity is also rising  
(embedded systems now programmed in JavaScript)
- Obvious privacy, security and vendor lock-in concerns

# Embedded problems

Devices have more power and run modern software  
yet they are really hard to  
develop, debug, maintain and extend

# Embedded problems

**Why?**

# Embedded problems

# Why?

Increased SW/HW complexity

+

Embedded-specific **resource constraints**

# Resource constraints

- Enough memory to run just a specific pre-built workload
- Cross-compiling and flashing/provisioning
- Special “Embedded Linux” distributions
- RT deadline requirements
- Ergonomics trade-offs, lack of HW ports
- Licensing requirements (no GPLv3...)
- Weird HW combinations, countless HW revisions
- Throw-away HW, planned obsolescence
- Low quality Out-Of-Tree drivers
- <Add your own pet-peave here>



# Creative solutions against constraints

- Debug symbol servers and remote GDB sessions
- Booting rootfs over the network
- Special protocols for diagnostics/log/trace
- Debug vs Release images, “developer mode”
- And so on



# Creative solutions against constraints

- Debug symbol servers and remote GDB sessions
- Booting rootfs over the network
- Special protocols for diagnostics/log/trace
- Debug vs Release images, “developer mode”
- And so on



**Here comes eBPF**

# Wait a minute

**Embedded-eBPF sounds like a solution in  
search of a problem...**

# Wait a minute

**Embedded-eBPF sounds like a solution in  
search of a problem...**

It kind of is.

“Embedded” engineers drooling over tools of “Cloud” engineers

Would like to have same system observability powers


Precedent: SMP on embedded





# Explaining eBPF / BCC in a few slides!

**BCC automates**



VM bytecode  
Kernel ↑ Userspace

Links at the end for better learning resources.



# VM running bytecode in the Linux kernel

**Bytecode loaded from userspace via bpf() syscall**

Verified for safety, unsafe => syscall rejects bytecode


**Bytecode compiled to native machine code**

**Native code inserted in execution paths**

Event-driven programming

**Native code runs and collects data**

**Data shared with userspace**





## How does userspace produce that bytecode?

0:	79 12 60 00 00 00 00 00 00	r2 = *(u64 *)(r1 + 96)
1:	7b 2a 98 ff 00 00 00 00 00	*(u64 *)(r10 - 104) = r2
2:	79 17 70 00 00 00 00 00 00	r7 = *(u64 *)(r1 + 112)
3:	85 00 00 00 0e 00 00 00 00	call 14
4:	bf 06 00 00 00 00 00 00 00	r6 = r0
5:	b7 09 00 00 00 00 00 00 00	r9 = 0
6:	7b 9a c0 ff 00 00 00 00 00	*(u64 *)(r10 - 64) = r9
7:	bf 73 00 00 00 00 00 00 00	r3 = r7
8:	07 03 00 00 18 00 00 00 00	r3 += 24
9:	bf a1 00 00 00 00 00 00 00	r1 = r10
11:	07 01 00 00 c0 ff ff ff	r1 += -64
12:	b7 02 00 00 08 00 00 00 00	r2 = 8
13:	85 00 00 00 04 00 00 00 00	call 4

## How does userspace produce that bytecode?

```
0: 79 12 60 00 00 00 00 00 00  
1: 7b 2a 98 ff 00 00 00 00 00  
2: 79 17 70 00 00 00 00 00 00  
3: 85 00 00 00 0e 00 00 00 00  
4: bf 06 00 00 00 00 00 00 00  
5: b7 09 00 00 00 00 00 00 00  
6: 7b 9a c0 ff 00 00 00 00 00  
7: bf 73 00 00 00 00 00 00 00  
8: 07 03 00 00 18 00 00 00 00  
9: bf a1 00 00 00 00 00 00 00  
11: 07 01 00 00 c0 ff ff ff  
12: b7 02 00 00 08 00 00 00 00  
13: 85 00 00 00 04 00 00 00 00
```

**Directly write it byte by byte!**



```
r2 = *(u64 *) (r1 + 96)  
*(u64 *) (r10 - 104) = r2  
r7 = *(u64 *) (r1 + 112)  
call 14  
r6 = r0  
r9 = 0  
*(u64 *) (r10 - 64) = r9  
r3 = r7  
r3 += 24  
r1 = r10  
r1 += -64  
r2 = 8  
call 4
```



Clang can translate “restricted C” into eBPF bytecode  
Much easier than assembling bytes like the 1960s

**Still hard to write userspace interaction**



Clang can translate “restricted C” into eBPF bytecode  
Much easier than assembling bytes like the 1960s

## **Still hard to write userspace interaction**

### **BCC: the BPF Compiler Collection**

Framework to ease writing userspace eBPF programs

Abstracts Clang and `sys_bpf()` interaction

“restricted C” compiled & loaded in kernel on-the-fly

Provides Python, Lua and Go bindings

Provides production ready BCC-tools

# BCC program

```
#!/usr/bin/env python
from bcc import BPF

csrc = """
#include <uapi/linux/ptrace.h>

int kprobe__do_sys_open(struct pt_regs *ctx)
{
    char file_name[256];
    bpf_probe_read(&file_name, sizeof(file_name), PT_REGS_PARM1(ctx));
    bpf_trace_printk(fmt, sizeof(fmt), file_name);
}
"""

b = BPF(text=csrc)
b.attach_kprobe(event="do_sys_open", fn_name="kprobe__do_sys_open")
while True:
    time.sleep(1)
```

# BCC program

```
#!/usr/bin/env python
from bcc import BPF

csrc = """
#include <uapi/linux/ptrace.h>

int kprobe__do_sys_open(struct pt_regs *ctx)
{
    char file_name[256];
    bpf_probe_read(&file_name, sizeof(file_name), PT_REGS_PARM1(ctx));
    bpf_trace_printk(fmt, sizeof(fmt), file_name);
}
"""


b = BPF(text=csrc)
b.attach_kprobe(event="do_sys_open", fn_name="kprobe__do_sys_open")
while True:
    time.sleep(1)
```

The diagram illustrates the workflow of a BCC program:

- Annotations for the C source code block:**
  - Compiled to bytecode
  - Loaded & runs in kernel
  - Collects data
  - Sends to userspace
- Annotations for the BPF object creation and attach block:**
  - Calls Clang to compile above code
  - Loads bytecode via bpf()



**New tools** developed for the book **BPF Performance Tools: Linux System and Application Observability** by Brendan Gregg (Addison Wesley, 2019), which also covers prior BPF tools



Real  
power  
comes  
with the  
**BCC**  
**tools**



# Executive summary eBPF benefits

- System-wide observability
- No crashes / hangs
- No performance degradations
- Real-time production workload analysis
- Can be always enabled (no special debug builds) \*
- Fully upstream kernel feature, active community
- Big collection of production-ready tools
- More than just observing a system
  - Packet filtering, hw offloading



## Executive summary eBPF benefits

- System-wide observability
- No crashes / hangs
- No performance degradations
- Real-time production workload analysis
- Can be always enabled (no special debug builds)
- Fully upstream kernel feature, active community
- Big collection of production-ready tools
- More than just observing a system
  - Packet filtering, hw offloading

**Convincing, yes?**

# eBPF meets embedded

general / embedded-specific problems

multiple approaches

project advantages / disadvantages

trade-offs, no silver bullet

# General problem: portability / cross-compilation

Poking “outside” from within the eBPF VM

- VM has generic 64 bit instructions/registers/pointers
- Difficulty accessing 32 bit kernel/user data structures
- VM is capable of 32 bit register subaddressing
- Pointer arithmetic hacks can access 32bit offset data  
**Very fragile, not portable**
- Better solution: **BPF Type Format** adds type info to compiled eBPF  
(part of **C.O.R.E.**)



# General problem: portability / cross-compilation

## Portable eBPF (**C**ompile **O**nce, **R**un **E**verywhere)

- Dream: run precompiled eBPF on any machine and expect it to work
- Slimmer version of BCC using BTF info, no Clang runtime compilation  
(structure offsets built in BTF sections, macro identifiers → BPF variables)
- Current runtime compilation uses version/config specific C headers
  - Backwards, not forwards compatible
  - Manually copying non-UAPI structures to “restricted C”
  - Big variation of Linux kernel configs → *header* structures
- Kernel >= 5.2 can remove header filesystem dependency (kinda unrelated)
- Work on-going

# General problem: Security and unprivileged eBPF

Running eBPF programs requires root / CAP\_SYS\_ADMIN

- eBPF code is assumed not malicious
- CAP\_BPF will be added to restrict attack surface
- Unprivileged eBPF unlikely to happen

Care must be taken when running eBPF code in production

- Don't run arbitrary eBPF supplied by untrusted users
- Use additional security mechanisms like verified boot

Awesome (as always) relevant LWN.net article and comments:

<https://lwn.net/Articles/796328/>



# Approach 1: Precompiled eBPF + custom userspace

## PRO:

Lightest footprint possible  
(few kb C program)

Kernel provides helper libbpf  
(useful starting point)

## CON:

Need to write from scratch  
Userspace sys\_bpf() interaction

Can get complex, hard to maintain  
No pre-existing community

Some examples provided by Linux kernel tree  
in samples/bpf/

# Approach 2: Use BCC directly

## PRO:

- Vanilla upstream BCC
- Full framework capabilities
- All BCC-tools available
- Well tested, good performance

## CON:


- Installs and links against Clang
- Depends on Python (bcc-tools)
- ~ 300 MB storage

**Will benefit from C.O.R.E., but will still require python**

Example project: Androdeb  
(Requires > 2GB storage)



# Approach 3: BPFd



**Project abandoned due to high maintenance cost**


## PRO:

- 100 kb bin + libc dependency
- Full framework capabilities
- All BCC-tools available

## CON:

- Hard to maintain BCC<>BPFd interaction
- Host + target + transport dependent architecture

# Approach 4: DSL compiler from scratch - Ply



```
ply 'kprobe:i2c_transfer { print(stack); }'
```

## PRO:

- 50 kb bin + libc dependency
- High level, AWK-inspired DSL
- Self-contained
- Easy to build & deploy

## CON:

- Lack of kernel/user interaction control
- Lack of BCC-tools diversity
- Under heavy development
- Ply binary is not portable

# Approach 5: Replace BCC Python userspace with Go

## PRO:

~2 mb static-compiled eBPF loader


Full control over kernel/user interaction

Good coverage of BCC API bindings

## CON:

BCC-tools need rewriting in Go :)

Not much documentation



Full execsnoop reimplementation:

<https://github.com/iovisor/gobpf/blob/master/examples/bcc/execsnoop/execsnoop.go>





## Ways forward

- C.O.R.E. needs to be as successful as possible
- With C.O.R.E. BCC will be more lightweight
- Gobpf can eliminate the Python dependency
- BPFd reached a dead end
- Ply is standalone, will continue its awesomeness
- eBPF on embedded is already useful today<
- Much work remaining



## Recommended learning resources:

- LWN.net eBPF articles <https://lwn.net/>
- Brendan Gregg's blog: <http://www.brendangregg.com/blog/>
- BPF Performance Tools: Linux System and Application Observability, by Brendan Gregg, published by Addison Wesley (2019)
- Collabora eBPF blog posts  
<https://www.collabora.com/news-and-blog/blog/2019/04/05/an-ebpf-overview-part-1-introduction/>
- Internet Search has wealth of information on eBPF



# Thank you!





# Embedded Linux Conference

## North America



# Embedded Linux Conference

## North America